


The AI control plane for banking ·
Preview

Connect any AI. Govern every answer. Across every channel.

Bridge is the layer between your bank's AI and your customers — plug in any model, run it across chat and voice, ground it in your data, check every answer, and see everything in one place.

[Book a demo →](#)[See the platform](#)

One platform between your AI and your customers — connect, govern, observe.



83%*

**Questions auto-
resolved**

* design target

30×

**Cheaper per
simple query**

up to — via smart
routing

100%

**Answers pass the
guard**

chat and voice alike

6

**Regulatory
regimes**

audit evidence,
automated

**Your AI shouldn't be a dozen disconnected tools.
Bridge is the one layer every answer passes through —
typed or spoken.**

THE PLATFORM

Everything between your model and your customer

One pipeline does the whole job — connect, ground, govern, deliver and observe.



Connect any AI

OpenAI, Anthropic, Azure or a local model. Swap providers anytime — nothing downstream changes.

model-agnostic



Every channel

WhatsApp, the app, web chat and the call center — all running through one governed pipeline.

omnichannel



Voice & transcription

Take phone calls, transcribe speech to text, then answer or route. Voice gets the exact same guard as chat.

speech → text



Grounded in your data

Answers cite your manuals, policies and regulations (RAG) — not the model's guesswork.

RAG · citations



Governed & audited

An uncertainty guard on every answer, PII masked before the model (LGPD), and an append-only audit trail (BCB) mapped to SR 11-7.

guard · LGPD · audit



See everything



Live metrics, per-stage traces of every request, and AI-visibility monitoring of how your brand shows up across the models.

metrics · visibility

SEE IT IN ACTION

The actual console, running today

Not a mockup — real screenshots of Bridge: live metrics, the customer pipeline, and your model & channel connections.



Real screenshots of the Bridge console today (Next.js + FastAPI), shown in its Portuguese demo mode. The orange demo banner is cropped; everything else is the live UI.

THE FLOW

How a request flows through Bridge

From a customer's question — typed or spoken — to a safe answer, in one pass.



Step 1

Customer reaches out

A message or phone call comes in — WhatsApp, app, web or the call center. Voice is transcribed to text.



Step 2

Bridge prepares it

Personal data is masked, the question is grounded in your documents, and it's routed to the right-sized model.



Step 3

The guard decides

Confident → reply. Unsure → ask again or hand off to a human. No guessing in front of a customer.



Step 4

Everything is logged

Each step becomes a reproducible, audit-ready record your risk team can file — automatically.

Under the hood: 12 audited stages — data quality, governance, cache, routing, retrieval, the guard and the audit trail — each one a real, tested module.

THE GAINS

Your gains — in numbers

The values your risk, operations and finance teams can point to.



100% gated

No answer reaches a customer ungated

Every low-confidence answer — chat or voice — is held back or escalated instead of bluffed. Your AI stays quiet when it isn't sure.



83% resolved*

Most questions answered without a human

Routine queries are handled instantly; only the genuinely hard ones reach your team — so headcount goes to the cases that need judgment. *design target



up to 30x cheaper

Stop overpaying per question

Easy questions go to cheap models, repeats are served from cache, and only the complex ones reach a premium model. You stop paying top price for "what's my balance?".



6 regimes · automatic

Audit-ready without the busywork

Every answer is filed as SR 11-7 and EU AI Act evidence — PII masked (LGPD), append-only trail (BCB). No folder of screenshots.

Plugs into your stack — nothing answers without passing the guard

Models: OpenAI

Anthropic

Azure OpenAI

Local (Ollama)

Channels: WhatsApp

App

Web

Voice / call center

Built on the open-source `lub` engine — 22 calibration methods, 732 tests, Apache 2.0.

BUILT IN THE OPEN

The engine is open source. Audit it, run it, build on it.

Bridge runs on `lub` — our Apache-2.0 calibration engine. No black box: every estimator and metric is in the repo, with the tests to prove it.

22

calibration methods

732

tests passing

93%

code coverage

Apache

2.0

license



rafaelmartinsalves / **llm-uncertainty-banking**

Turn LLM uncertainty into auditor-ready regulatory evidence — the calibration engine under Bridge.

 Python Apache-2.0 ★ Star

[View on GitHub →](#)

BOOK A DEMO

See Bridge run on your own data.

- ✓ A 30-minute walkthrough on your channels and models
- ✓ Your documents, your guard thresholds
- ✓ Runs on a local model — no API keys to start

Prefer email? [Write to us directly.](#)

Request a demo

Tell us where to reach you — we'll follow up within a day.

Full name

Work email

Bank / company

What do you want to use Bridge for?

Request a demo →

We'll only use this to contact you about Bridge. Preview — no API keys required to try it.



Platform

How it works

GitHub

Book a demo

Bridge is the AI control plane built on the open-source lub engine (Apache 2.0). It is currently in preview and runs on a local model out of the box. Nothing here is legal, regulatory or model-validation advice, and Bridge does not by itself establish compliance with any framework. Figures, connection states and confidence scores shown on this page are illustrative.